

Pro T-SQL 2008 Programmer's Guide



Michael Coles

Pro T-SQL 2008 Programmer's Guide

Copyright © 2008 by Michael Coles

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN-13 (pbk): 978-1-4302-1001-6

ISBN-13 (electronic): 978-1-4302-1002-3

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Java™ and all Java-based marks are trademarks or registered trademarks of Sun Microsystems, Inc., in the US and other countries. Apress, Inc., is not affiliated with Sun Microsystems, Inc., and this book was written without endorsement from Sun Microsystems, Inc.

Lead Editors: Jonathan Gennick, Tony Campbell

Technical Reviewer: Adam Machanic

Editorial Board: Clay Andres, Steve Anglin, Ewan Buckingham, Tony Campbell, Gary Cornell, Jonathan Gennick, Matthew Moodie, Joseph Ottinger, Jeffrey Pepper, Frank Pohlmann, Ben Renow-Clarke, Dominic Shakeshaft, Matt Wade, Tom Welsh

Project Manager: Kylie Johnston

Copy Editor: Damon Larson

Associate Production Director: Kari Brooks-Copony

Production Editor: Elizabeth Berry

Compositor: Lynn L'Heureux

Proofreaders: Linda Seifert, April Eddy

Indexer: Broccoli Information Management

Artist: Kinetic Publishing Services, LLC

Cover Designer: Kurt Krames

Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail orders-ny@springer-sbm.com, or visit <http://www.springeronline.com>.

For information on translations, please contact Apress directly at 2855 Telegraph Avenue, Suite 600, Berkeley, CA 94705. Phone 510-549-5930, fax 510-549-5939, e-mail info@apress.com, or visit <http://www.apress.com>.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at <http://www.apress.com/info/bulksales>.

The information in this book is distributed on an “as is” basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

The source code for this book is available to readers at <http://www.apress.com>.

For Devoné and Rebecca

Contents at a Glance

About the Author	xvii
About the Technical Reviewer	xix
Acknowledgments	xxi
Introduction	xxiii
■ CHAPTER 1 Foundations of T-SQL	1
■ CHAPTER 2 T-SQL 2008 New Features	23
■ CHAPTER 3 Tools of the Trade	61
■ CHAPTER 4 Procedural Code and CASE Expressions	81
■ CHAPTER 5 User-Defined Functions	117
■ CHAPTER 6 Stored Procedures	151
■ CHAPTER 7 Triggers	187
■ CHAPTER 8 Encryption	219
■ CHAPTER 9 Common Table Expressions and Windowing Functions	247
■ CHAPTER 10 Integrated Full-Text Search	273
■ CHAPTER 11 XML	299
■ CHAPTER 12 XQuery and XPath	341
■ CHAPTER 13 Catalog Views and Dynamic Management Views	387
■ CHAPTER 14 SQL CLR Programming	407
■ CHAPTER 15 .NET Client Programming	451
■ CHAPTER 16 Data Services	495
■ CHAPTER 17 New T-SQL Features	525
■ CHAPTER 18 Error Handling and Dynamic SQL	553

■ CHAPTER 19	Performance Tuning	573
■ APPENDIX A	Exercise Answers	603
■ APPENDIX B	XQuery Data Types	613
■ APPENDIX C	Glossary	619
■ APPENDIX D	SQLCMD Quick Reference	631
■ INDEX	639

Contents

About the Author	xvii
About the Technical Reviewer	xix
Acknowledgments	xxi
Introduction	xxiii
CHAPTER 1 Foundations of T-SQL	1
A Short History of T-SQL	1
Imperative vs. Declarative Languages	1
SQL Basics	3
Statements	3
Databases	5
Transaction Logs	6
Schemas	7
Tables	7
Views	9
Indexes	9
Stored Procedures	10
User-Defined Functions	10
SQL CLR Assemblies	10
Elements of Style	11
Whitespace	11
Naming Conventions	13
One Entry, One Exit	15
Defensive Coding	18
SQL-92 Syntax Outer Joins	18
The SELECT * Statement	19
Variable Initialization	20
Summary	20
CHAPTER 2 T-SQL 2008 New Features	23
Productivity Enhancements	23
The MERGE Statement	26

New Data Types	34
Date and Time Data Types	34
The hierarchyid Data Type	38
hierarchyid Methods	45
Spatial Data Types	47
Grouping Sets	55
Other New Features	58
Summary	59

■ CHAPTER 3 **Tools of the Trade** 61

SQL Server Management Studio	61
SSMS Editing Options	63
Context-Sensitive Help	64
Graphical Query Execution Plans	66
Project Management Features	66
The Object Explorer	68
The SQLCMD Utility	69
Business Intelligence Development Studio	71
SQL Profiler	73
SQL Server Integration Services	75
The Bulk Copy Program	76
SQL Server 2008 Books Online	76
The AdventureWorks Sample Database	78
Summary	78

■ CHAPTER 4 **Procedural Code and CASE Expressions** 81

Three-Valued Logic	81
Control-of-Flow Statements	83
The BEGIN and END Keywords	83
The IF...ELSE Statement	85
The WHILE, BREAK, and CONTINUE Statements	87
The GOTO Statement	88
The WAITFOR Statement	89
The RETURN Statement	90
The TRY...CATCH Statement	91
The CASE Expression	93
The Simple CASE Expression	94
The Searched CASE Expression	95
CASE and Pivot Tables	97
COALESCE and NULLIF	103

	Cursors	104
	Summary	114
CHAPTER 5	User-Defined Functions	117
	Scalar Functions	117
	Recursion in Scalar User-Defined Functions	120
	Procedural Code in User-Defined Functions	123
	Multistatement Table-Valued Functions	133
	Inline Table-Valued Functions	143
	Restrictions on User-Defined Functions	147
	Nondeterministic Functions	147
	State of the Database	148
	Summary	149
CHAPTER 6	Stored Procedures	151
	Introducing Stored Procedures	151
	Calling Stored Procedures	152
	Managing Stored Procedures	155
	Stored Procedures in Action	156
	Recursion in Stored Procedures	163
	Table-Valued Parameters	173
	Temporary Stored Procedures	176
	Recompilation and Caching	176
	Stored Procedure Statistics	176
	Parameter Sniffing	179
	Recompilation	182
	Summary	184
CHAPTER 7	Triggers	187
	DML Triggers	187
	When to Use DML Triggers	188
	Auditing with DML Triggers	192
	Nested and Recursive Triggers	199
	The UPDATE and COLUMNS_UPDATED Functions	200
	Triggers on Views	205
	DDL Triggers	208
	Logon Triggers	213
	Summary	217

CHAPTER 8	Encryption	219
	The Encryption Hierarchy	219
	Service Master Keys	220
	Database Master Keys	221
	Certificates	223
	Asymmetric Keys	229
	Symmetric Keys	233
	Encryption Without Keys	241
	Hashing Data	241
	Extensible Key Management	243
	Transparent Data Encryption	244
	Summary	246
 CHAPTER 9	 Common Table Expressions and Windowing Functions	 247
	Common Table Expressions	247
	Multiple Common Table Expressions	249
	Recursive Common Table Expressions	253
	Windowing Functions	259
	The ROW_NUMBER Function	259
	The RANK and DENSE_RANK Functions	262
	The NTILE Function	267
	Aggregate Functions and OVER	269
	Summary	271
 CHAPTER 10	 Integrated Full-Text Search	 273
	iFTS Architecture	273
	Creating Full-Text Catalogs and Indexes	275
	Creating Full-Text Catalogs	275
	Creating Full-Text Indexes	277
	Full-Text Querying	282
	The FREETEXT Predicate	282
	The CONTAINS Predicate	285
	The FREETEXTTABLE and CONTAINSTABLE Functions	289
	Thesauruses and Stoplists	292
	Stored Procedures and Dynamic Management Views and Functions	296
	Summary	297

CHAPTER 11	XML	299
	Legacy XML	299
	OPENXML	299
	OPENXML Result Formats	303
	FOR XML Clause	307
	FOR XML RAW	307
	FOR XML AUTO	309
	FOR XML EXPLICIT	311
	FOR XML PATH	313
	The xml Data Type	315
	Untyped xml	316
	Typed xml	317
	The xml Data Type Methods	319
	The query Method	320
	The value Method	321
	The exist Method	322
	The nodes Method	323
	The modify Method	325
	XML Indexes	327
	XSL Transformations	332
	Summary	338
CHAPTER 12	XQuery and XPath	341
	XPath and FOR XML PATH	341
	XPath Attributes	343
	Columns Without Names and Wildcards	344
	Element Grouping	345
	The data Function	346
	XPath and NULL	348
	The WITH XMLNAMESPACES Clause	349
	Node Tests	350
	XQuery and the xml Data Type	351
	Expressions and Sequences	352
	The query Method	354
	Location Paths	355
	Node Tests	357
	Namespaces	359
	Axis Specifiers	361

Dynamic XML Construction	363
XQuery Comments	366
Data Types	366
Predicates	367
Conditional Expressions (if...then...else)	373
Arithmetic Expressions	374
XQuery Functions	375
Constructors and Casting	378
FLWOR Expressions	380
Summary	385
CHAPTER 13 Catalog Views and Dynamic Management Views	387
Catalog Views	387
Table and Column Metadata	388
Index Metadata	390
Querying Permissions	392
Dynamic Management Views and Functions	395
Session Information	396
Connection Information	398
Currently Executing SQL	398
Tempdb Space	399
Server Resources	400
Unused Indexes	400
INFORMATION_SCHEMA Views	402
Summary	404
CHAPTER 14 SQL CLR Programming	407
The Old Way	407
The SQL CLR Way	408
SQL CLR Assemblies	409
User-Defined Functions	412
Stored Procedures	419
User-Defined Aggregates	423
Creating a Simple UDA	424
Creating an Advanced UDA	428
SQL CLR User-Defined Types	435
Summary	448

CHAPTER 15	.NET Client Programming	451
	ADO.NET	451
	The .NET SQL Client	453
	Connected Data Access	453
	Disconnected Datasets	458
	Parameterized Queries	460
	Nonquery, Scalar, and XML Querying	465
	SqlBulkCopy	468
	Multiple Active Result Sets	475
	LINQ to SQL	480
	Using the O/RM Designer	480
	Querying with LINQ to SQL	482
	Inspecting the O/RM Classes	491
	Summary	493
CHAPTER 16	Data Services	495
	Introducing HTTP Endpoints	495
	Consuming HTTP Endpoints	499
	Web Services	503
	ADO.NET Data Services	509
	Creating an ADO.NET Data Service	510
	Creating an ADO.NET Data Service Consumer	517
	Summary	522
CHAPTER 17	New T-SQL Features	525
	Set Operators	525
	The OUTPUT Clause	527
	The TOP Keyword	531
	CROSS APPLY and OUTER APPLY	533
	The TABLESAMPLE Clause	535
	The NEWSEQUENTIALID Function	537
	Date and Time Functions	538
	The max Data Types	541
	Synonyms	543


FILESTREAM Support	544
Enabling FILESTREAM Support	545
Creating FILESTREAM Filegroups	546
FILESTREAM-Enabling Tables	548
Accessing FILESTREAM Data	549
Summary	550
 ■ CHAPTER 18 Error Handling and Dynamic SQL	553
Error Handling	553
Legacy Error Handling	553
Try...Catch Exception Handling	555
The RAISERROR Statement	557
Debugging Tools	558
PRINT Statement Debugging	559
Trace Flags	559
SSMS Integrated Debugger	560
Visual Studio T-SQL Debugger	561
Dynamic SQL	564
The EXECUTE Statement	564
SQL Injection and Dynamic SQL	564
Troubleshooting Dynamic SQL	567
The sp_executesql Stored Procedure	568
Dynamic SQL and Scope	568
Client-Side Parameterization	570
Summary	570
 ■ CHAPTER 19 Performance Tuning	573
SQL Server Storage	573
Files and Filegroups	573
Space Allocation	574
Data Compression	579
Indexes	590
Heaps	590
Clustered Indexes	591
Nonclustered Indexes	592
Filtered Indexes	596
Optimizing Queries	596
Reading Query Plans	597
Methodology	600
Summary	601

APPENDIX A	Exercise Answers	603
	Chapter 1	603
	Chapter 2	603
	Chapter 3	604
	Chapter 4	605
	Chapter 5	606
	Chapter 6	607
	Chapter 7	607
	Chapter 8	608
	Chapter 9	608
	Chapter 10	609
	Chapter 11	609
	Chapter 12	609
	Chapter 13	610
	Chapter 14	610
	Chapter 15	611
	Chapter 16	611
	Chapter 17	611
	Chapter 18	612
	Chapter 19	612
APPENDIX B	XQuery Data Types	613
APPENDIX C	Glossary	619
APPENDIX D	SQLCMD Quick Reference	631
	Command-Line Options	631
	Scripting Variables	635
	Commands	636
INDEX		639

About the Author

■ **MICHAEL COLES** is a Microsoft MVP with over a dozen years' experience in SQL database design, T-SQL development, and client-server application programming. He has consulted in a wide range of industries, including the insurance, financial, retail, and manufacturing sectors, among others. Michael's specialty is developing and performance-tuning high-profile SQL Server-based database solutions. He currently works as a consultant for a business intelligence consulting firm. He holds a degree in information technology and multiple Microsoft and other certifications.

Michael has published dozens of technical articles online and in print magazines, including SQL Server Central, ASPToday, and *SQL Server Standard*. Michael is the author of the books *Pro T-SQL 2005 Programmer's Guide* (Apress, 2007) and *Pro SQL Server 2008 XML* (Apress, 2008), and he contributed to *Accelerated SQL Server 2008* (Apress, 2008). His current projects include coauthoring the book *Pro SQL Server 2008 Full-Text Search*.



About the Technical Reviewer

■ **ADAM MACHANIC** is a Boston-based independent database consultant, writer, and speaker. He has been involved in dozens of SQL Server implementations for both high-availability OLTP and large-scale data warehouse applications, and has optimized data access layer performance for several data-intensive applications. Adam has written for numerous web sites and magazines, including SQLblog, Simple Talk, Search SQL Server, SQL Server Professional, CoDe, and VSJ. He has also contributed to several books on SQL Server, including *Expert SQL Server 2005 Development* (Apress, 2007) and *Inside SQL Server 2005: Query Tuning and Optimization* (Microsoft Press, 2007). Adam regularly speaks at user groups, community events, and conferences on a variety of SQL Server and .NET-related topics. He is a Microsoft Most Valuable Professional (MVP) for SQL Server and a Microsoft Certified IT Professional (MCITP).



Acknowledgments

I've said it before, and I'll say it again—delivering books like this into your hands takes the coordinated efforts of dozens of people working toward a common goal. There's no way you would be reading these words right now if not for the entire team at Apress. This book is the product of the work of all my Apress teammates.

With that in mind, I would like to start by thanking my editors Tony Campbell and Jonathan Gennick, who pulled this project together and oversaw it from the first page of the dedication to the last page of the index. I would also like to thank the hardest-working project manager ever, Kylie Johnston, who kept everyone on track and on schedule. I want to send a special thank you to my technical reviewer (and one of my favorite SQL authors), Adam Machanic, for keeping me honest and challenging me to “explain, explain, explain!” at every opportunity. I would also like to thank Elizabeth Berry, Damon Larson, Linda Seifert, and April Eddy, the team members who ensured that I expressed myself as clearly and cleanly as possible during the copy edit and proofreading phases. I also want to thank the numerous other team members who spent countless hours laying out pages, manipulating images, and contributing their skills in a variety of ways to bring this book to you.

I would like to thank my family, including my girlfriend, Donna; my mom; Eric; Jennifer; Chris; Desmond; and Deja. I'd also like to thank my aunt Linda and her family for their support.

Most important, thank you to Devoné and Rebecca—my little angels—for keeping a smile on my face.

Finally, I would like to send a special thank you to Steve Jones, Microsoft MVP and entrepreneur extraordinaire of SQL Server Central fame, and Chuck Heinzelman, also an MVP and editor-in-chief of *SQL Server Standard* magazine, for their continued support.

Introduction

I still recall the first “database” application I ever wrote. It was a Turbo Pascal–based application for state government, designed to keep an inventory of tools and hazardous waste materials for a state institution in the late 1980s. I recall running into a lot of issues, including performance, large data storage, extensibility, and data integrity. I mention this only because these are just the types of problems that modern enterprise-class SQL DBMSs are specifically designed to handle. What’s more, they abstract away the internal workings (well, most of them anyway) so that you can concentrate more on your data and less on writing code to manipulate it. As an example, a simple sort algorithm that consumed over 100 lines of code in my custom Turbo Pascal “data-base” application is whittled down to a single `ORDER BY` clause in SQL.

This abstraction allows you to spend less time worrying about *how things get done* and more time thinking about *what you want to get done*. Although I don’t use the term in everyday conversation, I can say that this change in thinking about storage represents a true “paradigm shift.” The new version of SQL Server builds on the foundation laid out by previous releases, adding new capabilities and functionality designed to meet the increasing demands of a sophisticated developer base.

This book was originally scheduled to be an update of my *Pro T-SQL 2005 Programmer’s Guide* book. The sheer number of new features, however, demanded a nearly complete rewrite. I designed this new book with the goal of helping T-SQL developers get the absolute most out of the exciting new development features and functionality in SQL Server 2008.

Who This Book Is For

This book is intended for SQL Server developers who need to port code from prior versions of SQL Server, and those who want to get the most out of T-SQL on the 2008 release. You should have a working knowledge of SQL, preferably T-SQL on SQL Server 2005 or 2000, as most of the examples in this book are written in T-SQL. In this book, I will cover some of the basics of T-SQL, including some introductory concepts like data domain and three-valued logic—but this is not a beginner’s book. I will not be discussing database design, database architecture, normalization, and the most basic of SQL constructs in any kind of detail. Instead I will be focusing most of my discussion on topics of new SQL Server 2008 functionality, which assumes a basic understanding of SQL statements like `INSERT` and `SELECT`.

A working knowledge of C# and the .NET Framework is also useful (but not required), as some examples in the book will be written in C#. When C# sample code is provided, it is explained in detail, so an in-depth knowledge of the .NET Framework class library is not required.

How This Book Is Structured

This book was written to address the needs of three types of readers:

- SQL developers who are coming from other platforms to SQL Server 2008
- SQL developers who are moving from prior versions of SQL Server to SQL Server 2008
- DBAs and nondevelopers who need a working knowledge of T-SQL functionality to effectively support SQL Server 2008 instances

For all types of readers, this book is designed to act as a tutorial that describes and demonstrates new T-SQL features with working examples, and as a reference for quickly locating details about specific features. The following sections provide a chapter-by-chapter overview.

Chapter 1

Chapter 1 starts this book off by putting SQL Server 2008's implementation of T-SQL in context, including a short history of T-SQL, a discussion of T-SQL basics, and an overview of T-SQL coding best practices.

Chapter 2

Chapter 2 dives right into the new features of T-SQL on SQL Server 2008, with a discussion of productivity-enhancing features, the new MERGE statement, new data types like geometry and hierarchyid, and grouping sets.

Chapter 3

Chapter 3 gives an overview of the newest generation of tools available to SQL Server developers. Tools discussed include SQL Server Management Studio (SSMS), SQLCMD, Business Intelligence Development Studio (BIDS), and SQL Profiler, among others.

Chapter 4

Chapter 4 introduces T-SQL procedural code, including control-of-flow statements like IF...THEN and WHILE. In this chapter, I also discuss CASE expressions and CASE-derived functions, and provide an in-depth discussion of SQL three-valued logic.

Chapter 5

Chapter 5 discusses the various types of T-SQL user-defined functions available to encapsulate T-SQL logic on the server. I talk about all forms of T-SQL-based user-defined functions, including scalar user-defined functions, inline table-valued functions, and multistatement table-valued functions.

Chapter 6

Chapter 6 covers stored procedures, which allow you to create server-side T-SQL subroutines. In addition to describing how to create and execute stored procedures on SQL Server, I also address a thorny issue for some—the issue of why you might want to use stored procedures.

Chapter 7

Chapter 7 introduces all three types of SQL Server triggers: classic DML triggers, which fire in response to DML statements; DDL triggers, which fire in response to server and database DDL events; and logon triggers, which fire in response to server LOGON events.

Chapter 8

Chapter 8 discusses SQL Server encryption functionality, including the column-level encryption functionality introduced in SQL Server 2005 and the new transparent database encryption (TDE) and extensible key management (EKM) functionality, both introduced in SQL Server 2008.

Chapter 9

Chapter 9 dives into the details of common table expressions (CTEs) and windowing functions in SQL Server 2008, which feature the `OVER` clause.

Chapter 10

Chapter 10 discusses the advancements made to SQL Server 2008 integrated full-text search (iFTS), including greater integration with the SQL Server query engine and greater transparency by way of new iFTS-specific data management views and functions.

Chapter 11

Chapter 11 provides an in-depth discussion of SQL Server 2008 XML functionality, which carries forward the new features introduced in SQL Server 2005 and improves upon them. I cover several XML-related topics in this chapter, including the `xml` data type and its built-in methods, the `FOR XML` clause, and XML indexes.

Chapter 12

Chapter 12 discusses XQuery and XPath support in SQL Server 2008. SQL Server 2008 improves on the XQuery support introduced in SQL Server 2005, including support for the `xml` data type in XML DML insert statements and the `let` clause in FLWOR expressions.

Chapter 13

Chapter 13 introduces SQL Server 2008 catalog views, which are the preferred tools for retrieving database and database object metadata. This chapter also discusses dynamic management views and functions, which provide access to server and database state information.

Chapter 14

Chapter 14 is a discussion of SQL Common Language Runtime (SQL CLR) functionality in SQL Server 2008. In this chapter, I discuss and provide examples of SQL CLR stored procedures, user-defined functions, user-defined types, and user-defined aggregates. I also talk about the restrictions that have been removed in SQL CLR support in SQL Server 2008.

Chapter 15

Chapter 15 focuses on client-side support for SQL Server, including ADO.NET-based connectivity and one of Microsoft's newest technology offerings, LINQ to SQL.

Chapter 16

Chapter 16 discusses SQL Server connectivity using middle-tier technologies. Since native HTTP endpoints are deprecated in SQL Server 2008, I discuss them as items that may need to be supported in existing databases but should not be used for new development. I focus instead on possible replacement technologies, such as ADO.NET Data Services and IIS/.NET Web Services.

Chapter 17

Chapter 17 switches the focus back to T-SQL with a discussion of additional SQL Server 2008 features that were carried forward from their initial introduction in SQL Server 2005; features like the INTERSECT and EXCEPT operators, the DML statement OUTPUT clause, and improvements to the TOP clause. I also discuss some additional features and functionality that are new in SQL Server 2008, like new date and time functions and FILESTREAM support.

Chapter 18

Chapter 18 discusses improvements to server-side error handling made possible with the TRY...CATCH block. I also discuss various methods for debugging code, including using the Visual Studio T-SQL debugger. This chapter wraps up with a discussion of dynamic SQL and SQL injection, including the causes of SQL injection and methods you can use to protect your code against this type of attack.

Chapter 19

Chapter 19 provides an overview of performance-tuning SQL Server code. This chapter discusses SQL Server storage, indexing mechanisms, and query plans. I wrap up the chapter with a discussion of my own personal methodology for troubleshooting T-SQL performance issues.

Appendix A

Appendix A provides the answers to the exercise questions that I've included at the end of each chapter.

Appendix B

Appendix B is designed as a quick reference to the XQuery Data Model (XDM) type system.

Appendix C

Appendix C provides a quick reference glossary to several terms, many of which may be new to those using SQL Server for the first time.

Appendix D

Appendix D is a quick reference to the SQLCMD command-line tool, which allows you to execute ad hoc T-SQL statements and batches interactively, or run script files.

Conventions

To help make reading this book a more enjoyable experience, and to help you get as much out of it as possible, I've used the following standardized formatting conventions throughout.

C# code is shown in code font. Note that C# code is case sensitive. Here's an example:

```
while (i < 10)
```

T-SQL source code is also shown in code font, with keywords capitalized. Note that I've lowercased the data types in the T-SQL code to help improve readability. Here's an example:

```
DECLARE @x xml;
```

XML code is shown in code font with attribute and element content in bold for readability. Some code samples and results have been reformatted in the book for easier reading. XML ignores whitespace, so the significant content of the XML has not been altered. Here's an example:

```
<book publisher = "Apress">Pro SQL Server 2008 XML</book>:
```

■ **Note** Notes, tips, and warnings are displayed like this, in a special font with solid bars placed over and under the content.

SIDEBARS

Sidebars include additional information relevant to the current discussion and other interesting facts. Sidebars are shown on a gray background.

Prerequisites

This book requires an installation of SQL Server 2008 to run the T-SQL sample code provided. Note that the code in this book has been specifically designed to take advantage of new SQL Server 2008 features, and most of the code samples will not run on prior versions of SQL Server. The code samples presented in the book are designed to be run against the AdventureWorks 2008 sample database, available from the CodePlex web site at www.codeplex.com/MSFTDBProdSamples.

If you are interested in compiling and deploying the .NET code samples (the client code and SQL CLR examples) presented in the book, I highly recommend an installation of Visual Studio 2008. Although you can compile and deploy .NET code from the command line, I've provided instructions for doing so through the Visual Studio Integrated Development Environment (IDE). I find that the IDE provides a much more enjoyable experience.

Some examples, such as the ADO.NET Data Services examples in Chapter 16, require an installation of IIS as well. Other code samples presented in the book may have specific requirements, such as the LINQ samples, which require the .NET Framework 3.5. I've added notes to code samples that have additional requirements like these.

Downloading the Code

The sample code for this book is available in a ZIP file in the Downloads section of the Apress web site at www.apress.com. The ZIP file is structured so that each subdirectory contains all the sample code for its corresponding chapter.

Contacting the Author

The Apress team and I have made every effort to ensure that this book is free from errors and defects. Unfortunately, the occasional error does slip past us, despite our best efforts. In the event that you find an error in the book, please let us know! You can submit errors to Apress by visiting www.apress.com, locating the book page for this book, and clicking Submit Errata. Alternatively, feel free to drop a line directly to me at michaelco@optonline.net.